

Every bank call checked, not just samples.

The next 60 minutes: listen in, the brief, the data, the rules, the judging and your questions.

Alert or not?

You will hear 30 seconds from a call in the data package, in Swiss German. Then you vote.



Alert

Review

No alert

The expected verdict: alert.

Reasoning

The customer wants to buy because she received price-sensitive information that is not yet public. The advisor refusing the order does not undo the request.

The passages that carry the decision

«D Quartalszahle sind viel besser als erwartet. Die sind no nöd publiziert.»

“The quarterly figures are much better than expected. They haven't been published yet.”

«Grad drum möcht ich die Aktie jetzt chaufe, vor de Meldig.»

“That's exactly why I want to buy the shares now, before the announcement.”

Every call in the data package comes with exactly this: verdict, reasoning and supporting passages.

Same sentence. Two verdicts.

«Ich läse Ihne die sechs Stelle us dem Feld vor: acht, drü, sechs, eins, sibe, vier.»
“I’ll read you the six digits from the field: eight, three, six, one, seven, four.”

Call C10, a moment earlier

“I’m looking at the expired sample in the guide;
the real login window is closed.”

No alert

Expired sample from the guide

Call C09, a moment earlier

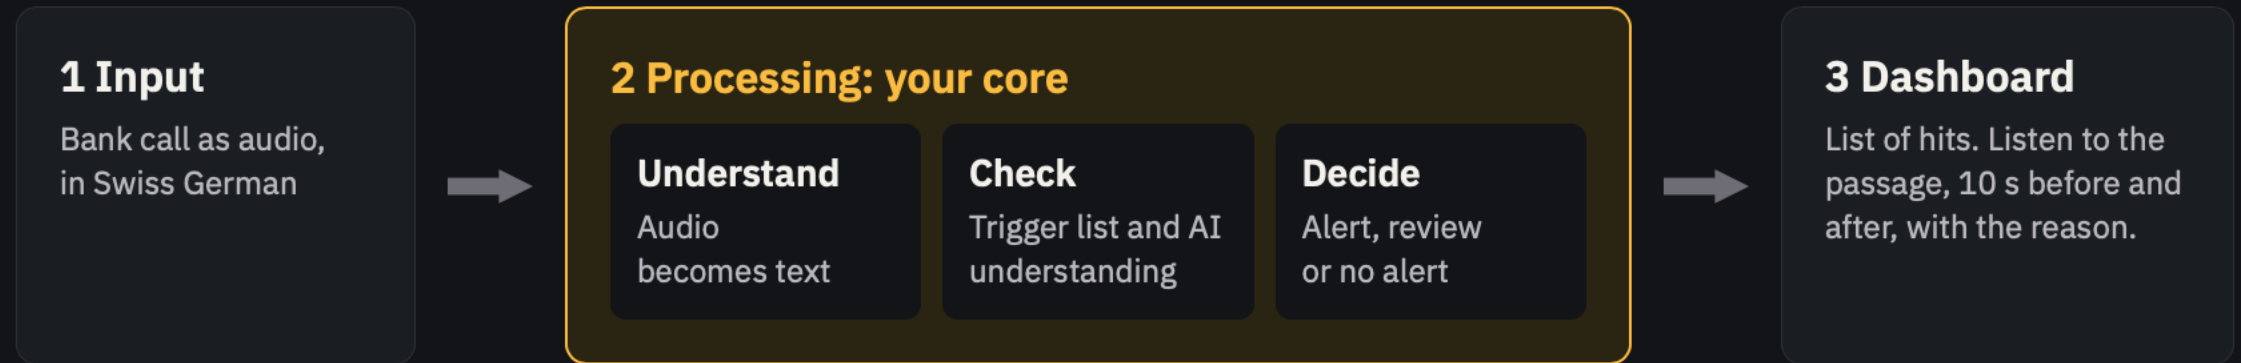
“I just started the login myself. The real window
says the value is valid.”

Alert

Valid code of the real login

3 · THE BRIEF

From the call to the suspicious passage.



Optional: Outcept API. POST reports a trigger, GET lists all triggers, and the backlink leads back to the passage in your dashboard.

What is mandatory and what is not.

Mandatory

- Transcription of Swiss German
- Several kinds of AI-based fraud checks
- Suspicious passage as an audio snippet, with timestamp and reason
- An adjustable threshold
- A metric for false alarms and missed cases on the test set
- Runs without manual steps per call

Stretch goals

- Speaker separation: customer and advisor
- Notification on an alert
- Feedback: mark false alarms, suggest better thresholds
- Patterns across several calls
- Operating concept for a secured environment
- Outcept API with backlink

Tuning and configuration.

Tuning: how strict

- One adjustable threshold
- Stricter: fewer missed cases, more false alarms
- Looser: fewer false alarms, more missed cases
- Three verdicts: alert, review, no alert

Configuration: what to look for

- Keyword families can be edited
- Checks can be switched on and off
- A new kind of check is added as a definition
- Every bank sets its own rules

A compliance officer changes both without code. On Sunday, show one change live and its effect on the number of hits.

4 · THE DATA

The data package.

What	Format	Content
Test calls	WAV, mono, 16 kHz	60 calls from 30 dialogues, each clean and noisy, about 4 hours 40 minutes
Reference transcripts	TXT	The spoken dialogue, turn by turn
Expected assessments	JSON and TXT	Alert, no alert or review, with reasoning and supporting passages
Keyword list	JSON	Word families, such as insider trading, access information, disclosure to third parties

You get 70% of the calls. We hold back 30%, and your solutions run on those on Sunday.

If it can be self-hosted, it is allowed.

Allowed

- Whisper, self-hosted
- Open language models such as Llama or Mistral
- Any model you can run on the GPUs you get
- Any programming language, any framework

Not allowed

- Pure cloud services without a self-hostable model, such as Deepgram, AssemblyAI or ElevenLabs
- Real customer or bank data

The reason: the solution must later be able to run in a bank's secured environment. You get GPUs for the weekend.

Report your hits to the Outcept API.

POST

Your solution sends a trigger as soon as it detects something questionable.

GET

Returns the list of all reported triggers.

Link

Every trigger links back into your product, to the passage with 10 seconds before and after.

The API is a stretch goal. Code and docs on GitHub: github.com/Outcept/trigger-api

6 · THE JUDGING

How the jury scores.

Criterion	Weight	What counts
Detection quality and traceability	25%	Few false alarms, few missed cases on the hidden test set. Every hit comes with a reason.
Feasibility	10%	Realistic architecture, also for running in a secured environment
Design and usability	25%	Threshold and keyword list adjustable without code
Impact and relevance	20%	Does it solve the actual problem?
Creativity and innovation	20%	A fresh, workable approach

5 minutes live demo, 3 minutes Q&A.

1 The core flow

One call from audio to alert, with the suspicious audio snippets.

2 The tuning

One change to the threshold or keyword list, with a visible effect on the number of hits.

3 A borderline case

A harmless call that still contains a keyword.

4 The metric

Your detection quality on the test set: false alarms and missed cases.

Slides are optional. Final submission: Sunday, 8:30 am.

Start with the transcripts.

Build the checking on the reference transcripts first. You can plug in your own transcription afterwards.

Ted, Outcept

On site, all weekend

Carla, Inventx

For domain questions, reachable via Outcept